

DS 642: Applications of Parallel Computing

Shahriar Afkhami


Week 6-7:

Introduction to CUDA and Graphics Processors

What is GPU?


- Programming with GPU follows vector processors in that all processors do the same task simultaneously. In that sense, they offer substantial speed-up over sequential programs, or can over-perform programs written for distributed memory machines.
- Before GPU was introduced, CPU did all the graphic processing tasks.
- In the early 1990s, computer manufactures began to include GPU into computer systems with the aim of accelerating common graphics routines.
- GPU manufactures, such as NVIDIA and ATI/AMD, found a way to use GPUs for more general purposes, not just for graphics or videos.

What is the difference between GPU and CPU?





- The major difference between GPU and CPU is that GPU has a highly parallel structure.
- GPU's are more effective than CPU's if used on data that can be partitioned and processed in parallel.
- GPU is highly optimized to perform advanced calculations such as floating-point arithmetic, matrix arithmetic and so on.

Heterogeneous Parallel Computing





- Different goals produce different designs
 - Manycore assumes work load is highly parallel
 - Multicore must be good at everything, parallel or not
- Multicore: minimize latency experienced by 1 thread
 - lots of big on-chip caches
 - extremely sophisticated control
- Manycore: maximize throughput of all threads
 - lots of big ALUs
 - multithreading can hide latency ... so skip the big caches
 - simpler control, cost amortized over ALUs via SIMD

Reminder: SIMD



Operation count:
4 loads, 4 multiplies, and 4 saves



Operation count:
1 loads, 1 multiplies, and 1 saves

- Single Instruction Multiple Data architectures make use of data parallelism
- Amortize control overhead over SIMD width
 - OpenMP / Pthreads / MPI all neglect SIMD parallelism
 - Because it is difficult for a compiler to exploit SIMD
 - Many languages (like C) are difficult to vectorize

Reminder: what's in a CPU?

Five-stage pipelined architecture:

- Instruction Fetch
- Instruction Decode
- Execution (ALU)
- Memory Access
- Write Back

To reduce latency, modern architectures have

- Out-of-order control logic
- Fancy branch predictor
- Memory pre-fetcher
- Large data cache

Hardware Big Ideas

- Remove components that make a single instruction stream run fast: Out-of-order, branch prediction, etc take space and energy.
- A larger number of (smaller simpler) cores: focus on data parallel (same operation, different data)
- Simplify branches: process branches sequentially and use masking to turn off thread (hardware handles all details of masking).
- Use threads to hide high latency operations (memory operations are 100s-1000s of cycles):
 - Remove fancy caches and pre-fetch logic that help avoid stalls (a core cannot run the next instruction because of a dependency on a previous one)
 - When one warp stalls, run another warp


Adding ALUs (Removing Decoders)

- Single Instruction Multiple Thread (SIMT)
 - NVIDIA calls this *warp* of threads
 - Warp operates with 32 threads (logical SIMD width)
 - Share one fetch/decode across warp of threads
 - Each thread execs same instructions at once
 - Some may be inactive (for conditional exec)
- Close to SIMD: single instruction multiple data
- Organize threads into logical grids of blocks
- Several types of device memory

Basic GPU architecture (Nvidia)

- GPUs many processors are called Streaming Processors (SP); each streaming processor is capable of executing a sequential thread.

- Streaming Processors are grouped in a Streaming Multiprocessor (SM).



- Each multi-processor is equipped with warp scheduler: responsible for threads control, special function units: to handle transcendentals and double-precision operations, a set of registers, and configurable shared memory.

The CUDA Programming Model


- CUDA is a programming model designed for:
 - Many-core architectures
 - Wide SIMD parallelism
 - Scalability
- CUDA provides:
 - A thread abstraction to deal with SIMD
 - Synchronization & data sharing between small groups of threads
- CUDA programs are written in C++ with minimal extensions: set of extensions to enable heterogeneous programming
- OpenCL is inspired by CUDA, but HW & SW vendor neutral: similar programming model, C only for device code

Simple Processing Flow

```
memory_allocation();  
copy_input_from_host_to_device();  
load_and_execute_gpu_code<<<nBlk, nTid>>>(args);  
cudaDeviceSynchronize();  
copy_result_from_device_to_host();  
memory_cleanup();
```

- Copy input data from CPU memory to GPU memory
- Load GPU program and execute, caching data on chip for performance
- Copy results from GPU memory to CPU memory

Simple Processing Flow



- **Host:** The CPU and its memory (host memory)
- **Device:** The GPU and its memory (device memory)


CUDA “Hello World”

```
1  #include <stdio.h>
2
3  __global__ void hello(){
4
5      printf("Hello from block: %u, thread: %u\n",  );
6  }
7
8  int main(){
9
10     hello<<<  >>>();
11     cudaDeviceSynchronize();
12 }
```

Triple angle brackets mark a call to device code, also called kernel launch.


The parameters inside the triple angle brackets are the CUDA kernel execution configuration.

Hierarchy of Grids, Blocks, and Threads



- Parallel **kernels** composed of many threads: all threads execute the same sequential program
- Threads are grouped into **thread blocks**: threads in the same block can cooperate via extremely fast on-chip (but small) **shared memory**
- Threads/blocks have unique IDs
- Independent blocks can form grids

Hierarchy of Grids, Blocks, and Threads



- Registers per thread (private)
- Local cached memory per thread (lives in global memory!)
- Shared memory (shared in block)
- Global device level shared
- Constant cache shared by threads
 - Constant memory only written by host
- Texture cache shared by all block

Global memory 100 times slower than shared memory.

CUDA Function Declarations

`__global__`

- Kernel function (must return void)
- Invoked from host
- Executed in parallel on device


`__host__`

- Called and executed on host

`__device__`

- Called and executed on device
- Restrictions: no recursion, no static variables, etc.

CUDA thread/grid organization



Hierarchical Parallelism Strategy

- Use both blocks and threads.
- Limit on maximum number of threads per block (threads alone won't work for large arrays).
- Fast shared memory only between threads (blocks alone are slower).
- Implicit barrier between kernels

CUDA “Hello World”

```
1  #include <stdio.h>
2  #define NUM_BLOCKS 4
3  #define BLOCK_WIDTH 8
4
5  /* Function executed on device (GPU) */
6  __global__ void hello(void){
7
8      printf("Hello from GPU: thread %d and block %d\n",
9              threadIdx.x, blockIdx.x );
10 }
11 int main(void){
12
13     /* execute function on device (GPU) */
14     hello<<<NUM_BLOCKS, BLOCK_WIDTH>>>();
15
16     /* wait until all threads finish their job */
17     cudaDeviceSynchronize();
18 }
```

CUDA “Hello World”

- Threads (arranged in warps) form blocks
- Blocks (assigned to SMs) form grids
- Layout of blocks/threads can be up to 3D
- `threadIdx` struct gives thread ID
- `blockIdx` struct gives block ID
- Each struct has `x`, `y`, `z` fields
- `gridDim` gives blocks per grid
- `blockDim` gives threads per block (at most 1024 threads per block)

CUDA “Hello World”

Compile and build the program using NVIDIA's `nvcc` compiler:

```
nvcc -o helloCuda helloCuda.cu
```

`nvcc` separates source code into host and device components:


Device functions (e.g. `kernel()`) processed by NVIDIA compiler:

```
nvcc
```

Host functions (e.g. `main()`) processed by standard host compiler: `gcc`.

`nvcc` can compile both C and C++ codes, and the suffix should be `.cu`

Memory Model



Basic device memory management:

```
/* Explicit memory allocation returns pointers  
   to GPU memory */  
cudaMalloc()  
/* Explicit memory copy for host to/from device,  
   device to/from device */  
cudaMemcpy()  
/* Free memory */  
cudaFree()
```

First CUDA code

```
1  #include <stdio.h>
2  #define NUM_BLOCKS 4
3  #define BLOCK_WIDTH 8
4
5  /* Function executed on device (GPU) */
6  __global__ void hello(void){
7
8      printf("Hello from GPU: thread %d and block %d\n",
9              threadIdx.x, blockIdx.x );
10 }
11 int main(void){
12
13     /* execute function on device (GPU) */
14     hello<<<NUM_BLOCKS, BLOCK_WIDTH>>>();
15
16     /* wait until all threads finish their job */
17     cudaDeviceSynchronize();
18 }
```

Second CUDA code

A simple program that uses a GPU kernel function to add two numbers on the GPU, and then brings the results from the GPU back into the CPU for output. The Host Code:

```
1  int main(int argc, char* argv[]) {  
2  
3  // Read input values  
4      int a = atoi(argv[1]);  
5      int b = atoi(argv[2]);  
6  
7  // c is storage place for main memory result  
8      int c;  
9  
10 // dev_c is storage place for result on GPU (device)  
11     int *dev_c;  
12  
13 // Allocate memory on GPU to store result  
14     cudaMalloc((void*)&dev_c, sizeof(int));
```


Second CUDA code

```
1 // Use one GPU unit to perform addition and store result
2 // in dev_c
3     add<<<1,1>>>(a, b, dev_c);
4
5 // Move result into main memory from GPU
6     cudaMemcpy(&c, dev_c, sizeof(int), cudaMemcpyDeviceToHost);
7
8     printf("%d + %d = %d\n", a, b, c);
9
10 // Free the allocated memory on GPU.
11     cudaFree(dev_c);
12
13     return 0;
14 }
```

Second CUDA code

```
1 // Use one GPU unit to perform addition and store result
2 // in dev_c
3     add<<<1,1>>>(a, b, dev_c);
4
5 // Move result into main memory from GPU
6     cudaMemcpy(&c, dev_c, sizeof(int), cudaMemcpyDeviceToHost);
7
8     printf("%d + %d = %d\n", a, b, c);
9
10 // Free the allocated memory on GPU.
11     cudaFree(dev_c);
12
13     return 0;
14 }
```

Second CUDA code

Device code:

```
1 // Output is a pointer to address where result will be stored.
2 // The address is assumed to be a GPU memory address.
3
4 __global__ void add(int a, int b, int* c) {
5     *c = a + b;
6 }
```

Measuring Performance

The tool we use to measure the time GPU spends on a task is the CUDA Event API.

```
1 // define events in the field
2 cudaEvent_t      start, stop;
3 // create two events we need
4 HANDLE_ERROR( cudaEventCreate( &start ) );
5 HANDLE_ERROR( cudaEventCreate( &stop ) );
6 // instruct the runtime to record the event start.
7 HANDLE_ERROR( cudaEventRecord( start, 0 ) );
8 // get stop time, and display the timing results
9 HANDLE_ERROR( cudaEventRecord( stop, 0 ) );
10 HANDLE_ERROR( cudaEventSynchronize( stop ) );
11 float    elapsedTime;
12 HANDLE_ERROR( cudaEventElapsedTime( &elapsedTime,
13                                     start, stop ) );
14 printf( "Time to generate:  %3.1f ms\n", elapsedTime );
```

Memory allocation and copying

The command `cudaMalloc()`, similar to `malloc()` command in standard C, tells the CUDA runtime to allocate memory on the device (Memory of GPU):

```
1 // allocate memory on the GPU
2 HANDLE_ERROR( cudaMalloc( (void**)&dev_a, N * sizeof(int) ) );
```

We can use `cudaMemcpy()` from host code to place data in memory on a device, similar to `memcpy()` command in standard C:

```
1 // copy arrays 'a'
2 HANDLE_ERROR( cudaMemcpy( dev_a, a, N * sizeof(int),
3                          cudaMemcpyHostToDevice ) );
```

Or transfer data back from device (GPU) to the host (CPU)

```
1 cudaMemcpy( a, dev_a, N * sizeof(int),
2            cudaMemcpyDeviceToHost );
```

Example of vector addition

The Device Code:

```
1  
2  
3 __global__ void vector_add(float *out, float *a, float *b, int n) {  
4     for(int i = 0; i < n; i ++){  
5         out[i] = a[i] + b[i];  
6     }  
7 }
```

The Kernel Invocation:

```
1 vector_add<<<1,1>>>>(d_out, d_a, d_b, N);
```

This vector addition in CUDA uses only one GPU thread.

Example of vector addition

The Host Code:

```
1  int main() {
2      float *a, *b, *out;
3      float *d_a, *d_b, *d_out;
4
5      // Allocate host memory
6      a = (float*)malloc(sizeof(float) * N);
7      b = (float*)malloc(sizeof(float) * N);
8      out = (float*)malloc(sizeof(float) * N);
9
10     // Initialize host arrays
11     for(int i = 0; i < N; i++){
12         a[i] = 1.0f;
13         b[i] = 2.0f;
14     }
15
16     // Allocate device memory
17     cudaMalloc((void**)&d_a, sizeof(float) * N);
18     cudaMalloc((void**)&d_b, sizeof(float) * N);
19     cudaMalloc((void**)&d_out, sizeof(float) * N);
20
21     // Transfer data from host to device memory
22     cudaMemcpy(d_a, a, sizeof(float) * N, cudaMemcpyHostToDevice);
23     cudaMemcpy(d_b, b, sizeof(float) * N, cudaMemcpyHostToDevice);
24
25     // Executing kernel
26     vector_add<<<1,1>>>>(d_out, d_a, d_b, N);
27 }
```

Example of vector addition

The Host Code Continued:

```
1
2  // Transfer data back to host memory
3  cudaMemcpy(out, d_out, sizeof(float) * N, cudaMemcpyDeviceToHost);
4
5  // Verification
6  for(int i = 0; i < N; i++){
7      assert(fabs(out[i] - a[i] - b[i]) < MAX_ERR);
8  }
9  printf("out[0] = %f\n", out[0]);
10 printf("PASSED\n");
11
12 // Deallocate device memory
13 cudaFree(d_a);
14 cudaFree(d_b);
15 cudaFree(d_out);
16
17 // Deallocate host memory
18 free(a);
19 free(b);
20 free(out);
21 }
```


On Bridges-2, GPU nodes are compute capability 70 (Volta and Tesla):

```
1 $ nvcc add_vector.cu -o add_vector -lm \  
2 -arch=compute_70 -code=sm_70
```

which is a shorthand for:

```
1 -gencode=arch=compute_70,code=sm_70 \  
2 -gencode=arch=compute_70,code=compute_70
```

Examples of vector addition in parallel

The Device Code:


```
1
2 __global__ void vector_add_gird(float *out, float *a, float *b) {
3     out[threadIdx.x] = a[threadIdx.x] + b[threadIdx.x];
4 }
5
6 __global__ void vector_add_block(float *out, float *a, float *b) {
7     out[blockIdx.x] = a[blockIdx.x] + b[blockIdx.x];
8 }
9
10 __global__ void vector_add(float *out, float *a, float *b, int n) {
11     int index = threadIdx.x;
12     int stride = blockDim.x;
13
14     for(int i = index; i < n; i += stride){
15         out[i] = a[i] + b[i];
16     }
17 }
```

The Kernel Invocation:

```
1 vector_add<<<1,N>>>>(d_out, d_a, d_b);
2 vector_add<<<N,1>>>>(d_out, d_a, d_b);
3 vector_add<<<1,256>>>>(d_out, d_a, d_b, N);
```

Examples of vector addition in parallel


For the k -th thread, the loop starts from k -th element and iterates through the array with a loop of 256.



- `threadIdx.x` index of the thread within the block
- `blockDim.x` size of thread block (number of threads in the thread block)
- `blockIdx.x` index of the block within the grid
- `gridDim.x` size of the grid

Combining Blocks and Threads

To take advantage of CUDA GPUs, kernel should be launched with multiple thread blocks. We will use multiple thread blocks to create N threads; each thread processes an element of the arrays:




To assign a thread to a specific element, we need to know a unique index for each thread:

```
1 int index = threadIdx.x + blockIdx.x * blockDim.x;
```

Combining Blocks and Threads

```
1
2 __global__ void vector_add(float *out, float *a, float *b, int n) {
3     int index = blockIdx.x * blockDim.x + threadIdx.x;
4     out[index] = a[index] + b[index];
5
6 }
7
8 __global__ void vector_add(float *out, float *a, float *b, int n) {
9     int index = blockIdx.x * blockDim.x + threadIdx.x;
10
11     // Handling arbitrary vector size
12     if (index < n){
13         out[index] = a[index] + b[index];
14     }
15 }
```



```
int index = threadIdx.x + blockIdx.x * M;
           =      5      +      2      * 8;
           = 21;
```

Combining Blocks and Threads

```
1
2 __global__ void vector_add(float *out, float *a, float *b, int n) {
3     int index = blockIdx.x * blockDim.x + threadIdx.x;
4     out[index] = a[index] + b[index];
5
6 }
7
8 __global__ void vector_add(float *out, float *a, float *b, int n) {
9     int index = blockIdx.x * blockDim.x + threadIdx.x;
10
11     // Handling arbitrary vector size
12     if (index < n){
13         out[index] = a[index] + b[index];
14     }
15 }
```

```
1 vector_add<<<N/THREADS_PER_BLOCK, THREADS_PER_BLOCK>>>
2 (d_out, d_a, d_b, N);
3
4 int block_size = 256;
5 int grid_size = (N + block_size - 1)/block_size;
6 vector_add<<<grid_size, block_size>>>(d_out, d_a, d_b, N);
```

NVIDIA provides a command line profiler tool; which give a more insight information of CUDA program performance.

To profile our vector addition, use following command

```
1 $ nvprof ./vector_add
```

Summary

- Threads in block share faster memory (may synchronize with barriers).
- Blocks in kernel share slow device memory (coordinate by atomic memory operations).
- To gain efficiency use multiple thread blocks per processor

Summary

- Manycore processors provide useful parallelism
- CUDA makes parallel programming easy to use SIMD
- CUDA SIMD allows highly scalable algorithm design and implementation
- CUDA explicit parallelism:
 - Programmer's responsibility to schedule resources
 - Decompose algorithm into kernels
 - Decompose kernels into blocks
 - Decompose blocks into threads